Archives & Records Association UK & Ireland

AI preparedness guidelines for archivists

Written by:

Prof. Giovanni Colavizza, University of Copenhagen and University of Bologna
and
Prof. Lise Jaillant, Loughborough University (UK)

Preferred citation:
Colavizza, Giovanni, and Lise Jaillant. *AI Preparedness Guidelines for Archivists*. February 2026. Archives & Records Association (UK & Ireland).
https://www.archives.org.uk/ai-preparedness-guidelines-for-archivists

ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

KØBENHAVNS UNIVERSITET

Loughborough University

AI Ready: Practical Guidelines for Archivists

Artificial Intelligence (AI) is now a regular topic of conversation in archives. Managers and stakeholders are asking whether AI can speed up description, identify sensitive content, or provide new forms of access. This document offers practical guidance on how to prepare archival collections for AI in ways that remain true to archival principles and ethical commitments.

The key message is simple: **AI can support archival work, but only when collections are made "AI-ready"** through careful preparation, documentation, and governance. Automation is a constrained necessity, not a magic solution.

What do we mean by "AI-ready"?

In these guidelines, we distinguish between two broad types of AI that archives are likely to encounter:

- **Task-specific AI**
  Models trained to perform a well-defined action, such as:
  - classifying record types;
  - detecting names, places, or dates;
  - flagging personal or sensitive information.

- **Generative AI**
  Large language models and related tools that:
  - summarise records or aggregates;
  - propose draft descriptions or keywords;
  - answer questions in natural language based on collections.

A particularly relevant approach for archives is **Retrieval-Augmented Generation (RAG)**. RAG first retrieves relevant records or metadata from a well-prepared collection, and then asks a generative model to draft an answer or summary based only on this retrieved material. This reduces "hallucinations" and keeps outputs grounded in your holdings.

Regardless of the tool, AI will only be useful and trustworthy if the underlying collections are prepared along four main dimensions: **completeness, metadata and access, data types and formats, and application-specific metrics**.

Pillar 1 – Completeness and excluded data

Completeness asks how far your digital data reflects the underlying collection:

- Are all items digitised, or only a subset?
- Are certain periods, creators, or communities missing or under-represented?
- Have some records been deliberately excluded (e.g. for legal or privacy reasons)?

You do not need perfect completeness to use AI. However, you do need to **document the degree of completeness and the reasons for any exclusions**. This should be recorded in collection-level metadata or short documentation notes that AI tools can also consult.

Good practice includes:

- stating clearly whether the digital corpus is complete, partial, or sampled;
- explaining why some material is absent (e.g. not yet digitised, legal restrictions, loss, selection);
- flagging known biases (e.g. strong focus on certain communities).

For **generative AI**, this documentation helps models—and users—understand the scope and limits of what can be said based on the data. For **task-specific AI**, it informs how far we can generalise from any trained model.

Pillar 2 – Metadata and access

High-quality **metadata** is foundational for both conventional access and AI applications. For AI readiness, the emphasis is on:

- **Item-level metadata**
  Even minimal item-level data (date, creator, brief title) greatly improves retrieval and AI-assisted description.
- **Provenance and relationships**
  Information about where records come from, how they are organised (series, sub-series), and how they relate to each other must be preserved and made readable by machines and humans.
- **Narrative and discursive metadata**
  Generative AI can work effectively with unstructured text. Curatorial notes, interpretive essays, collection-level descriptions, and contextual statements are valuable input. These narratives:
  - enrich AI-generated summaries;
  - help AI reflect historical and cultural nuance;
  - surface issues of power, silences, and harm.
- **Access conditions and sensitivity**
  Record in structured form which items or series are open, restricted, or closed, and why. AI systems must respect these access rules.

In multilingual contexts, document the languages present in the collection and in the metadata. Modern models can often work across languages, but clear labelling helps.

Pillar 3 – Data types, formats, and file structures

AI needs to be able to read and retrieve the data efficiently. For archives, this means:

- **Preserve provenance, do not overwrite it**
  Original file names, formats, dates, and folder structures are part of the archival record. It is important not to “clean” or normalise these in place. Instead:

    - preserve original files and structures;
    - create well-structured *derivative* copies for AI use.
- **Aim for consistent formats in derivatives**
  Standardise, where feasible, on a small set of AI-friendly formats:
    - text: UTF-8 plain text, XML, or similar;
    - images: TIFF or JPEG;
    - audiovisual: consistent codecs and containers.
- **Use clear and predictable file and folder naming in derivatives**
  File paths should:
    - indicate series or collection membership;
    - link reliably back to the original record or reference code;
    - support programmatic retrieval (e.g. via APIs).

The goal is not to erase the complexity of born-digital collections, but to **add a stable, documented layer** that AI tools can work with reliably.

Pillar 4 – Application-specific metrics and evaluation

Every AI project needs a way to decide whether it is working. Instead of generic metrics, define **application-specific metrics** that match your goals and users.

Examples:
- For AI-drafted descriptions:
    - proportion of drafts accepted with minor edits;
    - staff time saved per record compared to manual description;
    - archivist ratings of accuracy and appropriateness.
- For sensitivity-flagging tools:
    - how many true positives (correctly flagged sensitive items);
    - how many false positives (unnecessary flags);
    - whether critical issues are systematically missed.
- For RAG-based access tools:
    - precision and recall for test queries;
    - user satisfaction and understanding;
    - transparency of links back to records and finding aids.

Plan a simple evaluation protocol in advance: what will be measured, how, and by whom. This will help answer funders and managers who ask whether AI is delivering value and managing risk.

AI-assisted workflows for archives

With these pillars in place, archives can start exploring concrete AI-assisted workflows. Three promising areas, all under human supervision, are:

1. **Draft descriptive metadata and subject enhancement**
    - Automatic extraction of candidate titles, dates, and short summaries from text.

    - Suggestion of topical keywords based on controlled vocabularies.
    - Grouping of records into candidate series or sub-series.

All outputs must be reviewed by archivists and treated as non-authoritative suggestions.

2. **Sensitivity review and harmful content detection**
    - Automated flagging of likely personal identifiers, confidential information, or protected categories.
    - Identification of racist, sexist, or otherwise harmful language.
    - Prioritisation of items needing human review before release.

These tools help archivists focus attention, but do not replace their legal or ethical judgement.

3. **Access, discovery, and interpretive interfaces**
    - RAG-based systems that answer user questions with references to specific records.
    - Natural-language summaries of complex series or record types.
    - Multilingual access where models can work across languages represented in the collections.

Such systems should be transparent about which collections they cover and link clearly back to the underlying records and metadata.

<u>A short checklist</u>

Before launching an AI project over archival collections, you should be able to say "yes" to most of the following:

- We have a clearly defined problem and use case.
- We understand how complete (or partial) the digital corpus is, and we have documented gaps and exclusions.
- We have at least basic, coherent item- and collection-level metadata, including provenance and access conditions.
- We have created or planned standardised, well-documented derivative data in consistent formats, without altering originals.
- We have identified application-specific metrics and a simple evaluation plan.
- We have clear human-in-the-loop workflows for reviewing and approving AI outputs.

If these conditions are not yet in place, the most productive next step is not to deploy another tool, but to invest in **AI data preparedness**. That work is not separate from archival practice; it is an extension of what archivists already do—documenting, contextualising, and caring for records—so that any use of AI supports, rather than undermines, the integrity and value of archives.

This document was created by Prof. Giovanni Colavizza, University of Copenhagen and University of Bologna and Prof. Lise Jaillant, Loughborough University (UK) as one of the outputs from the FLAME (AI For Libraries, Archives and Museums) project funded by the Archives and Records Association UK & Ireland (ARA).

These guidelines were published by in February 2026 under a CC licence.

Archives and Records Association UK & Ireland
www.archives.org.uk l ara@archives.org.uk l 01823 327077
The Keep, Creech Castle, Taunton, Somerset, TA1 2DX
Charity Registered in England and Wales Number 1041063 l Charity Registered in Scotland Number SC041671 l Registered Company Limited by Guarantee in England and Wales Number 2969472